

How to model $p(x)$

- $\log \hat{p}(x) = \text{CNN}(x)$

Infeasible

- Use the chain rule

Cool, but slow to generate

x_1	x_2	x_3
x_4	x_5	x_6
x_7	x_8	x_9

[Oord, Aaron van den, Nal Kalchbrenner, and Koray Kavukcuoglu.

"Pixel recurrent neural networks." (2016)]

$$p(x_1, \dots, x_d)$$

$$= p(x_1)p(x_2 \mid x_1) \dots p(x_d \mid x_1, \dots, x_{d-1})$$

$$p(x_k \mid x_1, \dots, x_{k-1}) = \text{RNN}(x_1, \dots, x_{k-1})$$