Motivation: Rating Prediction

- Input: An opinionated text document **d**
- Output: Discrete rating **r** $\in$ {**1, 2, …, k**}
- Using regular text categorization techniques
  - Doesn't consider the order and dependency of the categories
  - The features distinguishing r=2 from r=1 may be the same as those distinguishing r=k from r=k-1 (e.g., positive words generally suggest a higher rating)
- Solution: Add order to a classifier (e.g., ordinal logistic regression)