Bigram language model

So that's what we get for n = 2:

$$p(\mathbf{w}) = \cancel{p(w_1)} p(w_2|w_1) \ldots p(w_k|w_{k-1})$$

$$p(w_1|start)$$

It's normalized separately for each sequence length!

p(this) + p(that) = 1.0

p(this this) + p(this is) + ... + p(built built) = 1.0

...